

Quant III

Lab 12: Tree-based Method

Junlong Aaron Zhou

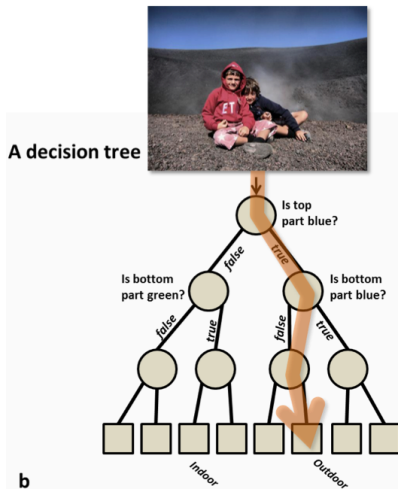
December 11, 2020

Announcement

- Final Exam:
 - I will distribute it Friday, December 18th, 9:45 am (**EST**)
 - Due: Sunday, December 20th, 11:59 pm (**EST**)
- Replication: due Friday, December 18th, 11:59 pm (**EST**)
- Final homework:
 - Due: Friday, December 18th, 11:59 pm (**EST**)
 - You should submit a pdf and a r object,
 - which includes a object for user: The input being a design matrix, and output being predictions.
- I add an extra office hour for next week, check **here**



A Binary Decision Tree

- binary tree: each node has either 2 children or 0 children





Binary Decision Tree on R^2

- Consider a binary tree on $\{(X_1, X_2) \mid X_1, X_2 \in R\}$



Binary Regression Tree on R^2

- Consider a binary tree on $\{(X_1, X_2) \mid X_1, X_2 \in R\}$



Recursive binary splitting

- ① Split the predictor space on one dimension (on predictor) into two sets.
- ② Split each of the resulting sets into two sets again.
- ③ Repeat until you have only a small number of observations left in each of the terminal sets.

Recursive binary splitting

- For a given j and s , split the space of each predictor x_j , $j = 1, \dots, p$, into two regions:

$$R_-(j, s) = \{\mathbf{x} | x_j < s\}$$

$$R_+(j, s) = \{\mathbf{x} | x_j \geq s\}$$

$R_-(j, s)$ is the region of predictor space where $x_j < s$.

- The splitting is done by finding j and s that minimizes

$$\sum_{i: \mathbf{x}_i \in R_-(j, s)} (y_i - \hat{y}_{R_-})^2 + \sum_{i: \mathbf{x}_i \in R_+(j, s)} (y_i - \hat{y}_{R_+})^2$$

where $\hat{y}_R = E(y | \mathbf{x}_i \in R)$ - Each resulting region is then split further until each region contains no more than N^* observations (where N^* is a small number).

Classification tree

- Suppose that the outcome variable is nominal $y_i \in \{c_1, \dots, c_K\}$ so that y_i can be one of K categories.
- Prediction for observation i : \hat{p}_{ik}
- Prediction in that region R_j : \hat{p}_{jk}
- The idea is exactly the same except that when calculating \hat{y}_{iR_m} we use the most frequent outcome category for observations in the set R_m .
- The only difference is that RSS is not used to find splits or to calculate the optimal complexity parameter α .

Impurity Measures

- Classification error:

$$1 - \max_k(\hat{p}_{jk})$$

where \hat{p}_{jk} is the proportion of observations in the region R_j that are from the class k . If a region R_j contains observations from many classes, then it does not classify them well. - Gini index


$$\sum_{k=1}^K \hat{p}_{jk}(1 - \hat{p}_{jk})$$

- Cross-entropy:

$$-\sum_{k=1}^K \hat{p}_{jk} \ln \hat{p}_{jk}$$

Two-Class Node Impurity Measures

- Consider binary classification
- Let p be the relative frequency of class 1.
- Here are three node impurity measures as a function of p



Bagging

- Combine predictions from multiple regression trees to improve their individual predictive capacity:
 - ① Take a bootstrap sample $b = 1, \dots, B$ from your data (with replacement).
 - ② Estimate a deep regression tree for sample b (without pruning).
 - ③ Calculate prediction $\hat{g}^b(\mathbf{x})$ given that tree.
 - ④ Calculate the *ensemble prediction*

$$\hat{g}(\mathbf{x}) = \sum_{b=1}^B \hat{g}^b(\mathbf{x}) / B$$

- Averaging over a large number of high-variance low-bias trees results in lower-variance bag of trees.

Random Forest

- Two disadvantages of bagging:
 - ① Computationally complex
 - ② Each tree are too similar
- Random forest solve these issues:
 - ① Take a bootstrap sample $b = 1, \dots, B$ from your data (with replacement).
 - ② Build regression tree by randomly selecting at each step $m < p$ predictors (usually $m \approx \sqrt{p}$).
 - ③ Calculate prediction $\hat{g}^b(\mathbf{x})$ from that tree.
 - ④ Calculate the *ensemble prediction*

$$\hat{g}(\mathbf{x}) = \sum_{b=1}^B \hat{g}^b(\mathbf{x}) / B$$

- Mathematically, RF should recover true $g(x)$

Different Forest

- You may see different forests, e.g.
 - Generalized Random Forest (Athey, et al.)
- Most the difference lies in the splitting criterion.

Optional: Boosting

- Ensemble methods combine multiple models
- **Parallel ensembles:**
 - each model is built independently
 - e.g. bagging and random forests
 - Main Idea: Combine many (high complexity, low bias) models to reduce variance
- **Sequential** ensembles:
 - Models are generated sequentially
 - Try to add new models that do well where previous models lack

Intuition

- Suppose we want to solve a classification problem (spam or not)
- A weak learner is a classifier that does slightly better than random.
 - Weak learners are like “rules of thumb”:
 - If an email has “Viagra” in it, more likely than not it’s spam.
 - Email from a friend is probably not spam.
 - A linear decision boundary.
- Can we combine a set of weak classifiers to form **single** classifier that makes accurate predictions
- Yes! Boosting solves this problem. [Rob Schapire (1990).]
- In general, you can use different classifier. (e.g., Grimmer et al, 2017)

Basic Idea

- Linear combination of basis functions:

$$f(x) = \sum_{m=1}^M v_m g_m(x)$$

- We'll consider learning by **empirical risk minimization**:

$$\hat{f} = \underset{f \in \mathbb{F}}{\operatorname{argmin}} \frac{1}{n} \sum_{i=1}^n \ell(y_i, f(x_i)),$$

for some **loss function** $\ell(y, \hat{y})$.

- Write ERM objective function as

$$J(v_1, \dots, v_M, h_1, \dots, h_M) = \frac{1}{n} \sum_{i=1}^n \ell \left(y_i, \sum_{m=1}^M v_m h_m(x) \right).$$

- How to optimize J ? i.e. how to learn?

Gradient-Based Methods

- **Suppose** our base hypothesis space is parameterized by $\Theta = R^b$:

$$J(v_1, \dots, v_M, \theta_1, \dots, \theta_M) = \frac{1}{n} \sum_{i=1}^n \ell \left(y_i, \sum_{m=1}^M v_m h(x; \theta_m) \right).$$

- Can we differentiate J w.r.t. v_m 's and θ_m 's? Optimize with SGD?
- For **some** hypothesis spaces and typical loss functions, yes!
- Neural networks fall into this category! (h_1, \dots, h_M are neurons of last hidden layer.)
- We have gradient boost machine, Adaboost, etc.
- For trees, we have gradient tree boosting